

PNEUMONIA_Detection_Using_X-Ray

April 9, 2024

```
[ ]: import warnings
warnings.filterwarnings('ignore')

[ ]: from tensorflow import keras

[ ]: from tensorflow.keras.layers import Input, Lambda, Dense, Flatten

[ ]: from keras.models import Model
from keras.applications.vgg16 import VGG16
from keras.applications.vgg16 import preprocess_input
from keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
import numpy as np
from glob import glob
import matplotlib.pyplot as plt

[ ]: IMAGE_SIZE = [224, 224]

train_path = 'Datasets/train'
valid_path = 'Datasets/test'

[ ]: vgg = VGG16(input_shape=IMAGE_SIZE + [3], weights='imagenet', include_top=False)

[ ]: for layer in vgg.layers:
    layer.trainable = False

[ ]: folders = glob('Datasets/train/*')
x = Flatten()(vgg.output)

[ ]: prediction = Dense(len(folders), activation='softmax')(x)
# create a model object
model = Model(inputs=vgg.input, outputs=prediction)
# view the structure of the model
model.summary()
```

Model: "functional_1"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 2)	50,178

Total params: 14,764,866 (56.32 MB)

Trainable params: 50,178 (196.01 KB)

Non-trainable params: 14,714,688 (56.13 MB)

```
[ ]: model.compile(  
    loss='categorical_crossentropy',  
    optimizer='adam',  
    metrics=['accuracy']  
)
```

```
[ ]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
[ ]: train_datagen = ImageDataGenerator(rescale = 1./255,  
                                       shear_range = 0.2,  
                                       zoom_range = 0.2,  
                                       horizontal_flip = True)  
  
test_datagen = ImageDataGenerator(rescale = 1./255)  
  
# Make sure you provide the same target size as initialied for the image size  
training_set = train_datagen.flow_from_directory('Datasets/train',  
                                                target_size = (224, 224),  
                                                batch_size = 10,  
                                                class_mode = 'categorical')  
  
test_set = test_datagen.flow_from_directory('Datasets/test',  
                                           target_size = (224, 224),  
                                           batch_size = 10,  
                                           class_mode = 'categorical')
```

Found 5216 images belonging to 2 classes.

Found 624 images belonging to 2 classes.

```
[ ]: r = model.fit(  
    training_set,  
    validation_data=test_set,  
    epochs=5,  
    steps_per_epoch=len(training_set),  
    validation_steps=len(test_set)
```

```
)
```

```
Epoch 1/5
522/522          359s 683ms/step -
accuracy: 0.8897 - loss: 0.2630 - val_accuracy: 0.8830 - val_loss: 0.4715
Epoch 2/5
522/522          0s 117us/step -
accuracy: 0.0000e+00 - loss: 0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss:
0.0000e+00
Epoch 3/5
522/522          362s 691ms/step -
accuracy: 0.9534 - loss: 0.1507 - val_accuracy: 0.8974 - val_loss: 0.4915
Epoch 4/5
522/522          0s 76us/step -
accuracy: 0.0000e+00 - loss: 0.0000e+00 - val_accuracy: 0.0000e+00 - val_loss:
0.0000e+00
Epoch 5/5
522/522          354s 676ms/step -
accuracy: 0.9617 - loss: 0.1237 - val_accuracy: 0.8478 - val_loss: 0.7690
```

```
[ ]: import tensorflow as tf
      from keras.models import load_model

      model.save('model_pneumonia.h5')
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```
[ ]: from keras.models import load_model
```

```
[ ]: from keras.preprocessing import image
```

```
[ ]: from keras.applications.vgg16 import preprocess_input
```

```
[ ]: import numpy as np
```


```
[ ]: model=load_model('model_pneumonia.h5')
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

```
[ ]: img=image.load_img('Datasets\\val\\NORMAL\\NORMAL2-IM-1427-0001.
      ↪jpeg',target_size=(224,224))
```

```
[ ]: x=image.img_to_array(img)
img
```

```
[ ]:
```



```
[ ]: x=np.expand_dims(x, axis=0)
x
```

```
[ ]: array([[[[ 6.,  6.,  6.],
               [ 0.,  0.,  0.],
               [ 7.,  7.,  7.],
               ...,
               [78., 78., 78.],
               [74., 74., 74.],
               [67., 67., 67.]],

            [[ 0.,  0.,  0.],
               [ 2.,  2.,  2.],
               [11., 11., 11.],
               ...,
               [82., 82., 82.],
               [69., 69., 69.],
               [64., 64., 64.]],

            [[ 0.,  0.,  0.],
               [ 5.,  5.,  5.],
               [12., 12., 12.],
               ...,
               [78., 78., 78.],
```

```

[70., 70., 70.],
[65., 65., 65.]],

...,

[[ 0.,  0.,  0.],
 [ 0.,  0.,  0.],
 [ 0.,  0.,  0.],
 ...,
 [ 0.,  0.,  0.],
 [ 0.,  0.,  0.],
 [ 0.,  0.,  0.]],

[[ 0.,  0.,  0.],
 [ 0.,  0.,  0.],
 [ 0.,  0.,  0.],
 ...,
 [ 0.,  0.,  0.],
 [ 0.,  0.,  0.],
 [ 0.,  0.,  0.]],

[[ 0.,  0.,  0.],
 [ 0.,  0.,  0.],
 [ 0.,  0.,  0.],
 ...,
 [ 0.,  0.,  0.],
 [ 0.,  0.,  0.],
 [ 0.,  0.,  0.]]], dtype=float32)

```

```

[ ]: img_data=preprocess_input(x)
img_data

```

```

[ ]: array([[[[ -97.939   , -110.779   , -117.68   ],
 [-103.939   , -116.779   , -123.68   ],
 [-96.939    , -109.779   , -116.68   ],
 ...,
 [-25.939003, -38.779    , -45.68    ],
 [-29.939003, -42.779    , -49.68    ],
 [-36.939003, -49.779    , -56.68    ]],

 [[-103.939   , -116.779   , -123.68   ],
 [-101.939   , -114.779   , -121.68   ],
 [-92.939    , -105.779   , -112.68   ],
 ...,
 [-21.939003, -34.779    , -41.68    ],
 [-34.939003, -47.779    , -54.68    ],
 [-39.939003, -52.779    , -59.68    ]]],

```

```

[[-103.939    , -116.779    , -123.68    ],
 [ -98.939    , -111.779    , -118.68    ],
 [ -91.939    , -104.779    , -111.68    ],
 ...,
 [ -25.939003,  -38.779    ,  -45.68    ],
 [ -33.939003,  -46.779    ,  -53.68    ],
 [ -38.939003,  -51.779    ,  -58.68    ]],

...,

[[-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ],
 ...,
 [-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ]],

[[-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ],
 ...,
 [-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ]],

[[-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ],
 ...,
 [-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ],
 [-103.939    , -116.779    , -123.68    ]]]], dtype=float32)

```

```
[ ]: classes=model.predict(img_data)
```

```
1/1          0s 297ms/step
```

```
[ ]: result=int(classes[0][0])
result
```

```
[ ]: 0
```

```
[ ]: if result==0:
      print("Person is Affected By PNEUMONIA")
    else:
      print("Person is not Affected By PNEUMONIA")
```

Person is Affected By PNEUMONIA

```
[ ]: import streamlit as st
```